

GSIP: More Questions

Frequently Asked, Stakeholder, Due Diligence, and Discussion Questions

Document version: 1.0 Last updated: February 2025

This document collects questions that stakeholders, investors, partners, and technical

reviewers of

Part A Product and Value

A.1 What is GSIP in one sentence?

A platform that turns natural-language questions into optimization results by running

real simulation

A.2 Who is it for?

Organizations that need to answer what if? and what's best? in complex domains (e.g.

finance, poli

A.3 Why not just use ChatGPT or another LLM?

General-purpose chatbots can invent plausible-looking statistics. GSIP uses AI only for

understandi

A.4 What makes the numbers trustworthy ?

Only domain-pack simulation code produces numeric results. Those results are persisted

immediately

A.5 What domains does it support today?

ToyPack (testing), FinancePack (portfolio backtesting), and SpatialPack (e.g.

pollution/diff

A.6 Can we add our own simulation logic?

Yes. You implement the domain pack contract (state schema, action schema, simulate,

score, feasibility

A.7 How long does a typical run take?

Depends on scenario count, fidelity mix, and hardware. With tens to low hundreds of

scenarios and

A.8 Can we use our own data (e.g. historical returns, maps)?

Yes where the domain pack supports it. State and actions can carry your parameters;

custom packs

Part B Technical

B.1 Where does the AI run?

AI (e.g. LLM) is used only in optional steps: objective formalization (parsing the

question) and

B.2 How is reproducibility guaranteed?

Deterministic scenario hashes (SHA-256 of inputs), fixed seeds per scenario, and stored

run spec and

B.3 What if the formalizer misinterprets the question?

You can extend keyword sets and prompts, or use the optional LLM step. Future UX may

allow editing

B.4 How do you scale to many scenarios?

Sim Fabric uses Ray for distributed execution; workers run scenarios in parallel. Caching

by scenario

B.5 What's in the run ledger?

PostgreSQL stores runs, scenarios, scenario_instances, metric_results, judge_scores,

artifacts (with

B.6 Is the system stateless?

No. Workflow state lives in Temporal; application state (runs, results) in PostgreSQL and

MinIO. Redi

B.7 How do you handle failures mid-run?

Temporal workflows and activities have retries and timeouts. Partial results are

persisted as

B.8 What about security and multi-tenancy?

JWT auth, RBAC (e.g. admin, analyst, viewer), tenant isolation at the data layer, and

audit events

Part C Business and Market

C.1 What s the business model?

Platform can be licensed (on-prem or private cloud), offered as SaaS (subscription or

usage-based

C.2 Who are the main competitors?

Generic AI assistants (no simulation guarantee); traditional simulation/optimization

tools (manua

C.3 What s the pricing model?

Not fixed in product; can be tiered by runs/scenarios/seats, with premium tiers for more

packs, SLA,

C.4 What s the go-to-market strategy?

Options: land in one vertical (e.g. finance or policy), design-partner pilots,

product-led

C.5 What s the revenue potential?

Depends on pricing, conversion, and market size. Architecture supports both SMB/team

(volume) and

C.6 Is it open source?

Core pipeline and SDK can be open (open-core model); specific domain packs or enterprise

features can

Part D Compliance and Risk

D.1 Can we use this in regulated industries (e.g. finance)?

Designed for audit: ledger, hashes, seeds, versioning, RBAC, audit events.

Sector-spec

D.2 How do you prevent garbage in, garbage out ?

Formalization quality is improved via keywords and optional LLM; rubrics and benchmarks

are versione

D.3 What if a domain pack has a bug?

Bugs affect only that pack's outcomes; they don't let an LLM invent numbers. Versioning

and run ledg

D.4 What about data residency and privacy?

Deployment can be on-prem or in a chosen cloud/region. Document ingestion (Evidence) and

any PII are c

D.5 How do you handle bias in objectives or rubrics?

Objectives and rubrics are explicit and versioned; bias is in the open. Best practice:

document ra

Part E Roadmap and Maturity

E.1 What's implemented today?

Core pipeline: API, Temporal workflow, formalization, scenario generation, Sim Fabric

(Ray), Judge

E.2 What's next (product)?

Optional problem-understanding and candidate-solution steps; more domain packs (e.g.

healthcare, I

E.3 What s next (operations)?

Production hardening: observability (traces, dashboards), runbooks, managed services

where appro

E.4 How do you decide quality done ?

Definition of Done and Quality Gates (reproducibility, evidence, simulation,

optimization

E.5 What are the main limitations today?

Keyword-based formalization can miss nuance; only three domain packs; no multi-user

collaboration

Part F Stakeholder and Due Diligence

F.1 Why will enterprises adopt this instead of building in-house?

Faster time-to-value (question report), built-in audit trail, and no need to build

formalization

F.2 What s the defensibility / moat?

Combination of: Non-Negotiable Truth architecture (trust), domain pack ecosystem

(vertical dep

F.3 What could kill the product?

Failure to close quality gates and deliver trust; slow domain pack rollout; poor UX;

strong simul

F.4 What s the team/skills need?

Product/engineering for platform and packs; domain experts for rubrics and benchmarks;

DevOps for

F.5 What would you do with more funding?

Accelerate domain packs and vertical depth; invest in formalization quality and UX (e.g.

scenario con

F.6 How do you measure success?

Product: run completion rate, scenarios per run, time to first result, reproducibility

pass rate. B

F.7 What s the exit or long-term vision?

Could be: standalone category leader in trustworthy AI-assisted decision support ;

acquisition b

F.8 What s the biggest open question?

Whether one or two verticals should be owned first (finance vs policy vs operations) and

whether go-t

Part G Discussion and Workshops

Use these for internal alignment or workshops:

1. Trust: What would it take for a compliance officer to approve GSIP for a regulated use

case in your

End of Questions Document